

# Parallel Computing

Ed Gonzalez

2022-12-01

I had to move onto the desktop version of R Studio for this as the cloud-based program did not allow me to use more than 0.5GB of RAM

```
install.packages("readr", repos = "http://cran.us.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/tz/sh20cj15711657_9_1d4v6m00000gn/T//Rtmpau3axm/downloaded_packages

library(readr)

airlines09 <- read_csv("/Users/ed/Downloads/2009.csv")

## Rows: 6429338 Columns: 28
## -- Column specification -----
## Delimiter: ","
## chr   (4): OP_CARRIER, ORIGIN, DEST, CANCELLATION_CODE
## dbl  (22): OP_CARRIER_FL_NUM, CRS_DEP_TIME, DEP_TIME, DEP_DELAY, TAXI_OUT, W...
## lgl   (1): Unnamed: 27
## date  (1): FL_DATE
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

Installing some packages for tidying the data set and for parallel computing

```
install.packages("tidyverse", repos = "http://cran.us.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/tz/sh20cj15711657_9_1d4v6m00000gn/T//Rtmpau3axm/downloaded_packages
```

```
install.packages("sparklyr", repos = "http://cran.us.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/tz/sh20cj15711657_9_1d4v6m00000gn/T//Rtmpau3axm/downloaded_packages

install.packages("multidplyr", repos = "http://cran.us.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/tz/sh20cj15711657_9_1d4v6m00000gn/T//Rtmpau3axm/downloaded_packages

library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.4.0      v dplyr 1.0.10
## v tibble 3.1.8       v stringr 1.5.0
## v tidyr 1.2.1        v forcats 0.5.2
## v purrr 0.3.5
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()

library(sparklyr)

##
## Attaching package: 'sparklyr'
##
## The following object is masked from 'package:purrr':
##
## invoke
##
## The following object is masked from 'package:stats':
##
## filter

library(multidplyr)
library(dplyr)
```

**Timing some code using dplyr that finds the mean arrival delay of each carrier**

```
system.time(
  airline_delay <- airlines09 %>%
  group_by(OP_CARRIER) %>%
  summarize(mean = mean(ARR_DELAY, na.rm = TRUE)) %>%
  print())
```

```
## # A tibble: 19 x 2
##   OP_CARRIER mean
##   <chr>      <dbl>
## 1 9E         0.950
## 2 AA         5.67
## 3 AS         1.32
## 4 B6         5.08
## 5 CO         5.51
## 6 DL         4.84
## 7 EV        11.7
## 8 F9         5.69
## 9 FL         8.20
## 10 HA        0.261
## 11 MQ         5.97
## 12 NW         3.92
## 13 OH        10.7
## 14 OO         3.30
## 15 UA         1.03
## 16 US         2.09
## 17 WN         1.66
## 18 XE         5.75
## 19 YV         6.04

##   user  system elapsed
## 0.127  0.031   0.159
```

Same as before, but now utilizing the for %do% process and timing it

```
library(foreach)
```

```
##
## Attaching package: 'foreach'

## The following objects are masked from 'package:purrr':
##
##   accumulate, when
```

```
carriers<-split(airlines09,airlines09$OP_CARRIER)
```

```
system.time(
  {foreach(i=carriers,.combine=c) %do%
    mean(i$ARR_DELAY, na.rm = T) %>%
    print()}
)
```

```
## [1] 0.9500134 5.6708979 1.3170298 5.0792806 5.5138186 4.8419736
## [7] 11.7207813 5.6856234 8.1978418 0.2607874 5.9679675 3.9242081
## [13] 10.6877569 3.3035038 1.0341901 2.0899596 1.6638597 5.7545910
## [19] 6.0439209
```

```
##      user  system elapsed
##    0.034   0.018   0.053
```

Using the `doParallel` package where I essentially just change the `%do%` to `%dopar%` and assign some cores

```
library(doParallel)
```

```
## Loading required package: iterators
```

```
## Loading required package: parallel
```

```
n.cores <- detectCores() - 1
registerDoParallel(n.cores)

system.time({
  foreach(i=carriers,.combine=c) %dopar%
    mean(i$ARR_DELAY, na.rm = T) %>%
    print()
})
```

```
## [1] 0.9500134 5.6708979 1.3170298 5.0792806 5.5138186 4.8419736
## [7] 11.7207813 5.6856234 8.1978418 0.2607874 5.9679675 3.9242081
## [13] 10.6877569 3.3035038 1.0341901 2.0899596 1.6638597 5.7545910
## [19] 6.0439209
```

```
##      user  system elapsed
##    0.053   0.088   0.043
```

Same procee utilizing the `multidplyr` package

```
library(multidplyr)
cl <- new_cluster(7)

system.time(
{
  airline_delay <- airlines09 %>%
    group_by(OP_CARRIER) %>%
    partition(cl) %>%
    summarize(mean = mean(ARR_DELAY, na.rm = TRUE)) %>%
    collect() %>%
    print()
})
```

```
## # A tibble: 19 x 2
##   OP_CARRIER    mean
```

```
##      <chr>      <dbl>
##  1 9E          0.950
##  2 NW          3.92
##  3 YV          6.04
##  4 AA          5.67
##  5 XE          5.75
##  6 AS          1.32
##  7 F9          5.69
##  8 HA          0.261
##  9 OH         10.7
## 10 WN          1.66
## 11 B6          5.08
## 12 FL          8.20
## 13 US          2.09
## 14 CO          5.51
## 15 MQ          5.97
## 16 DL          4.84
## 17 UA          1.03
## 18 EV         11.7
## 19 OO          3.30

##      user  system elapsed
##    3.399   1.997  17.210
```

## Creating a table with times for each computation

```
tidyverse.time <- c(0.128, 0.045, 0.233)
foreach.do.time <- c(0.032, 0.021, 0.052)
foreach.dopar.time <- c(0.004, 0.025, 0.085)
multidplyr.time <- c(3.73, 1.834, 7.336)

rbind(tidyverse.time, foreach.do.time, foreach.dopar.time, multidplyr.time)

##           [,1] [,2] [,3]
## tidyverse.time 0.128 0.045 0.233
## foreach.do.time 0.032 0.021 0.052
## foreach.dopar.time 0.004 0.025 0.085
## multidplyr.time 3.730 1.834 7.336
```

Using dplyr to select for specific variables and performing a summary function to find the mean of arrival delay, counting unique instances, and finding a linear regression model between arrival delay and taxi in/out times

```
system.time({
  taxi.times <- airlines09 %>%
```

```

select(ORIGIN, DEST, OP_CARRIER, ARR_DELAY, TAXI_IN, TAXI_OUT) %>%
drop_na() %>%
group_by(ORIGIN, DEST, OP_CARRIER) %>%
summarise(Count = n(), Mean = mean(ARR_DELAY, na.rm = T),
          fit = summary(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT))$r.squared)
print(head(taxi.times))
})

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

```

```
## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

## Warning in summary.lm(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT)): essentially perfect
## fit: summary may be unreliable

## 'summarise()' has grouped output by 'ORIGIN', 'DEST'. You can override using
## the '.groups' argument.

## # A tibble: 6 x 6
## # Groups:   ORIGIN, DEST [5]
##   ORIGIN DEST OP_CARRIER Count   Mean   fit
##   <chr>  <chr> <chr>      <int> <dbl> <dbl>
## 1 ABE    ATL    EV          766  17.0  0.116
## 2 ABE    CLE    XE          502  -5.89  0.201
## 3 ABE    CLT    US          354  -1.40  0.0790
## 4 ABE    CLT    YV           35   2.43  0.108
## 5 ABE    DTW    9E          969  -1.05  0.134
## 6 ABE    FLL    FL           62  10.2  0.0191

##   user  system elapsed
##  3.976   0.161   4.197
```

Doing the same as before but utilizing multidplyr functions to compare times

```
system.time({
  taxi.times2 <- airlines09 %>%
    select(ORIGIN, DEST, OP_CARRIER, ARR_DELAY, TAXI_IN, TAXI_OUT) %>%
```

```

drop_na() %>%
group_by(ORIGIN, DEST, OP_CARRIER) %>%
  partition(cl) %>%
  summarise(Count = dplyr::n(), Mean = mean(ARR_DELAY, na.rm = T),
            fit = summary(lm(ARR_DELAY ~ TAXI_IN + TAXI_OUT))$r.squared)
print(head(taxi.times2))
})

```

```

## # A tibble: 6 x 6
## # Groups:   ORIGIN, DEST [6]
##   ORIGIN DEST OP_CARRIER Count   Mean   fit
##   <chr> <chr> <chr>      <int> <dbl> <dbl>
## 1 ABE    ATL    EV          766 17.0  0.116
## 2 ABQ    DEN    F9         1042  3.28  0.0560
## 3 ABQ    HOU    WN         1000 -1.60  0.0258
## 4 ABQ    LBB    WN          363  2.15  0.0476
## 5 ABQ    MAF    WN          362  9.80  0.0183
## 6 ABQ    MDW    WN          724 -5.25  0.0312

```

```

##   user  system elapsed
##  1.589    0.566   14.289

```

## Using ggplot2 package to create a scatter plot

```

library(ggplot2)

scatter.p <- ggplot2::ggplot(taxi.times, aes(Count, fit))

scatter.p + geom_point() + theme_classic()


```

```

## Warning: Removed 2 rows containing missing values ('geom_point()').

```





# Creating boxplots for each individual airline

```
box.p <- ggplot2::ggplot(taxi.times, aes(OP_CARRIER, fit))
```

```
box.p+geom_boxplot(outlier.shape = NA ) + theme_classic()
```

## Warning: Removed 2 rows containing non-finite values ('stat\_boxplot()').

